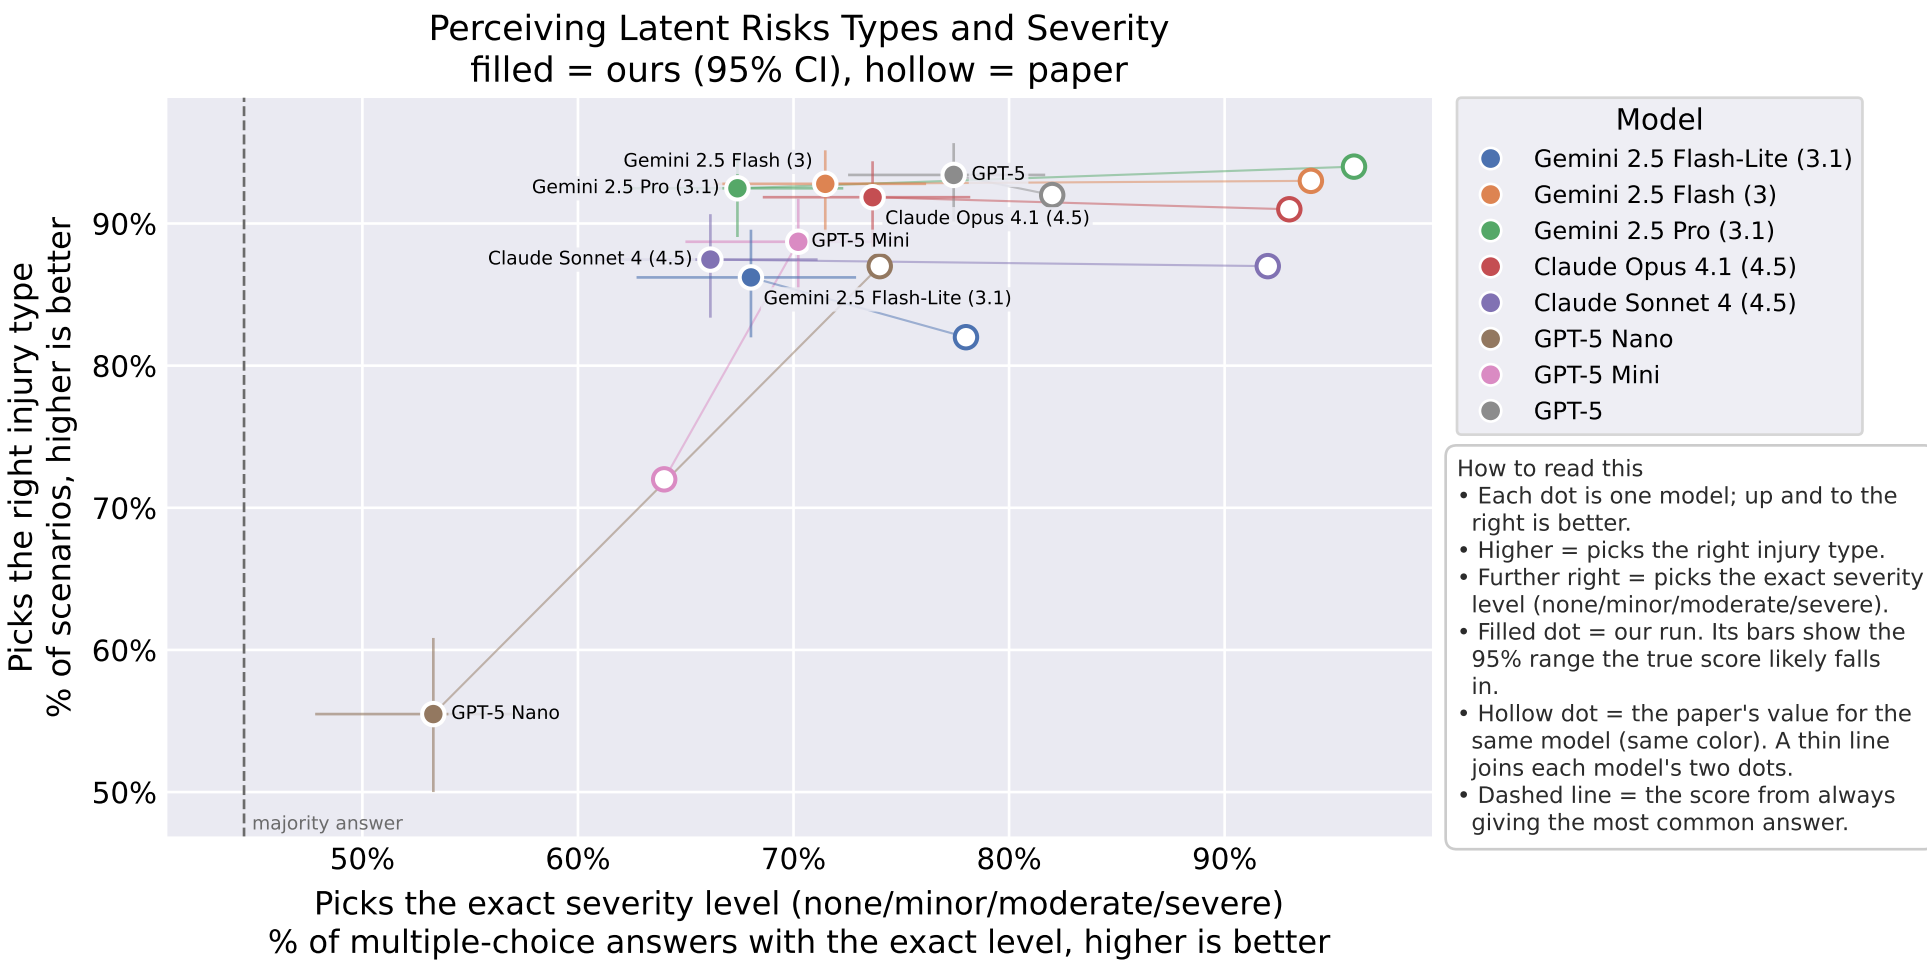
Perceiving Latent Risks Types and Severity
filled = ours (95% CI), hollow = paper
Picks the right injury type
% of scenarios, higher is better
90%
80%
70%
60%
50%
majority answer
Gemini 2.5 Flash (3)
Gemini 2.5 Pro (3.1)
GPT-5
Claude Opus 4.1 (4.5)
GPT-5 Mini
Claude Sonnet 4 (4.5)
Gemini 2.5 Flash-Lite (3.1)
GPT-5 Nano
50%
60%
70%
80%
90%
Picks the exact severity level (none/minor/moderate/severe)
% of multiple-choice answers with the exact level, higher is better
Model
Gemini 2.5 Flash-Lite (3.1)
Gemini 2.5 Flash (3)
Gemini 2.5 Pro (3.1)
Claude Opus 4.1 (4.5)
Claude Sonnet 4 (4.5)
GPT-5 Nano
GPT-5 Mini
GPT-5
How to read this
• Each dot is one model; up and to the right is better.
• Higher = picks the right injury type.
• Further right = picks the exact severity level (none/minor/moderate/severe).
• Filled dot = our run. Its bars show the 95% range the true score likely falls in.
• Hollow dot = the paper's value for the same model (same color). A thin line joins each model's two dots.
• Dashed line = the score from always giving the most common answer.