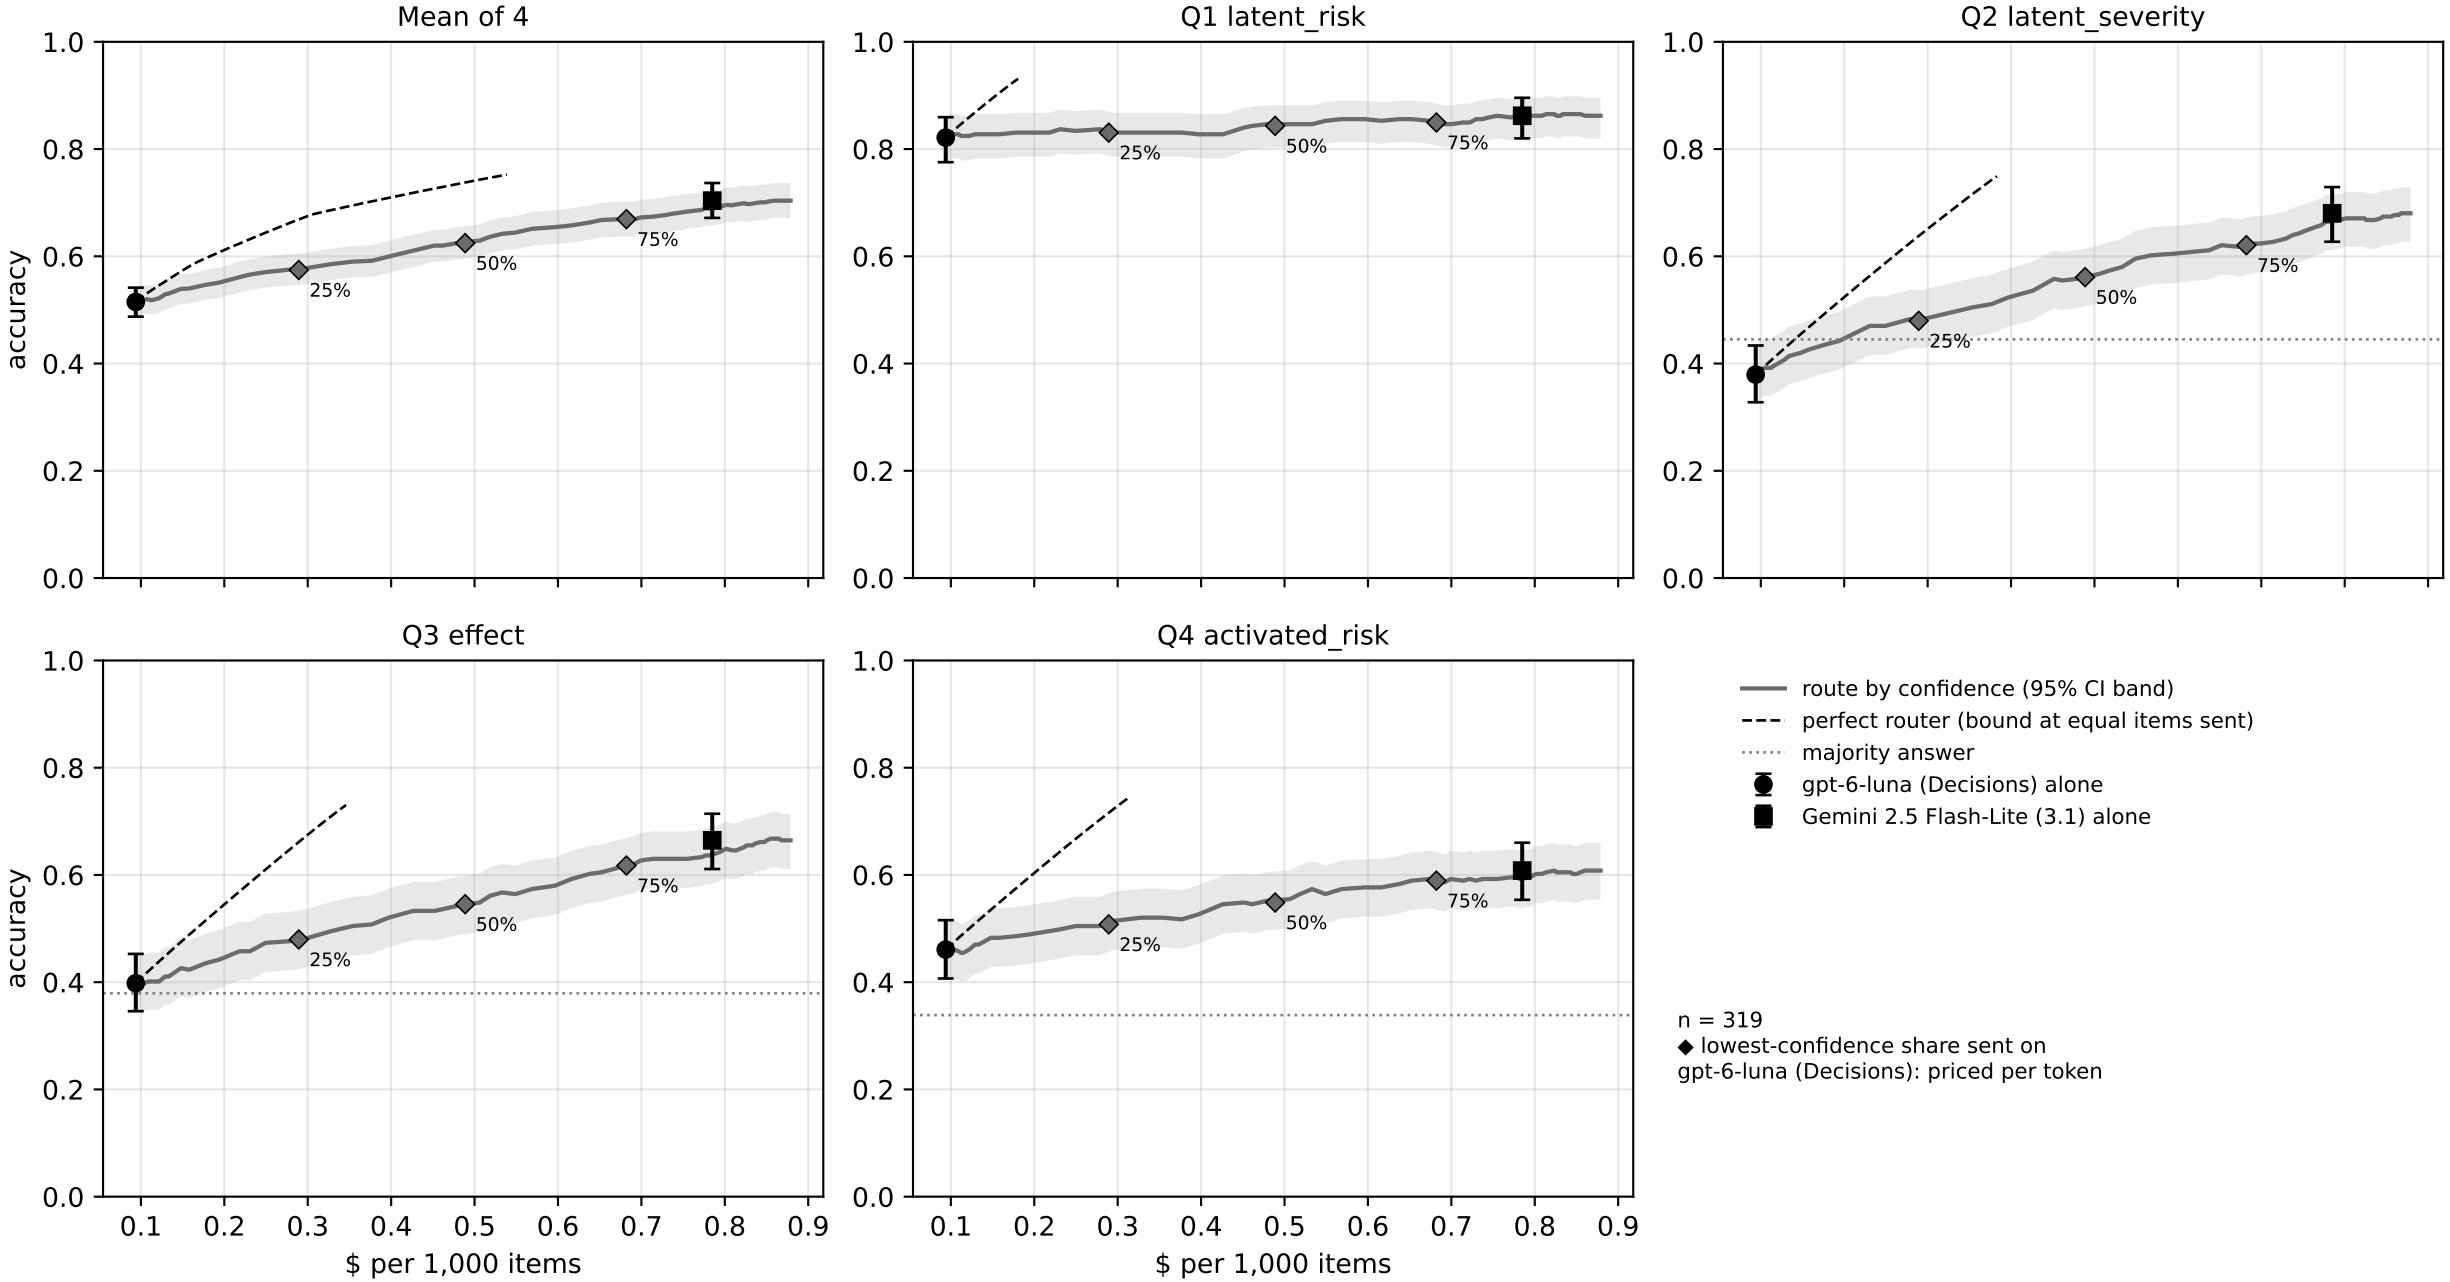
EXPLORATORY, post-hoc: escalation targets chosen after the Part 2 results were seen (not pre-registered)
gpt-6-luna (Decisions) first; low-confidence items to Gemini 2.5 Flash-Lite (3.1)
Mean of 4
Q1 latent_risk
Q2 latent_severity
Q3 effect
Q4 activated_risk
accuracy
$ per 1,000 items
25%
50%
75%
route by confidence (95% CI band)
perfect router (bound at equal items sent)
majority answer
gpt-6-luna (Decisions) alone
Gemini 2.5 Flash-Lite (3.1) alone
n = 319
◆ lowest-confidence share sent on
gpt-6-luna (Decisions): priced per token