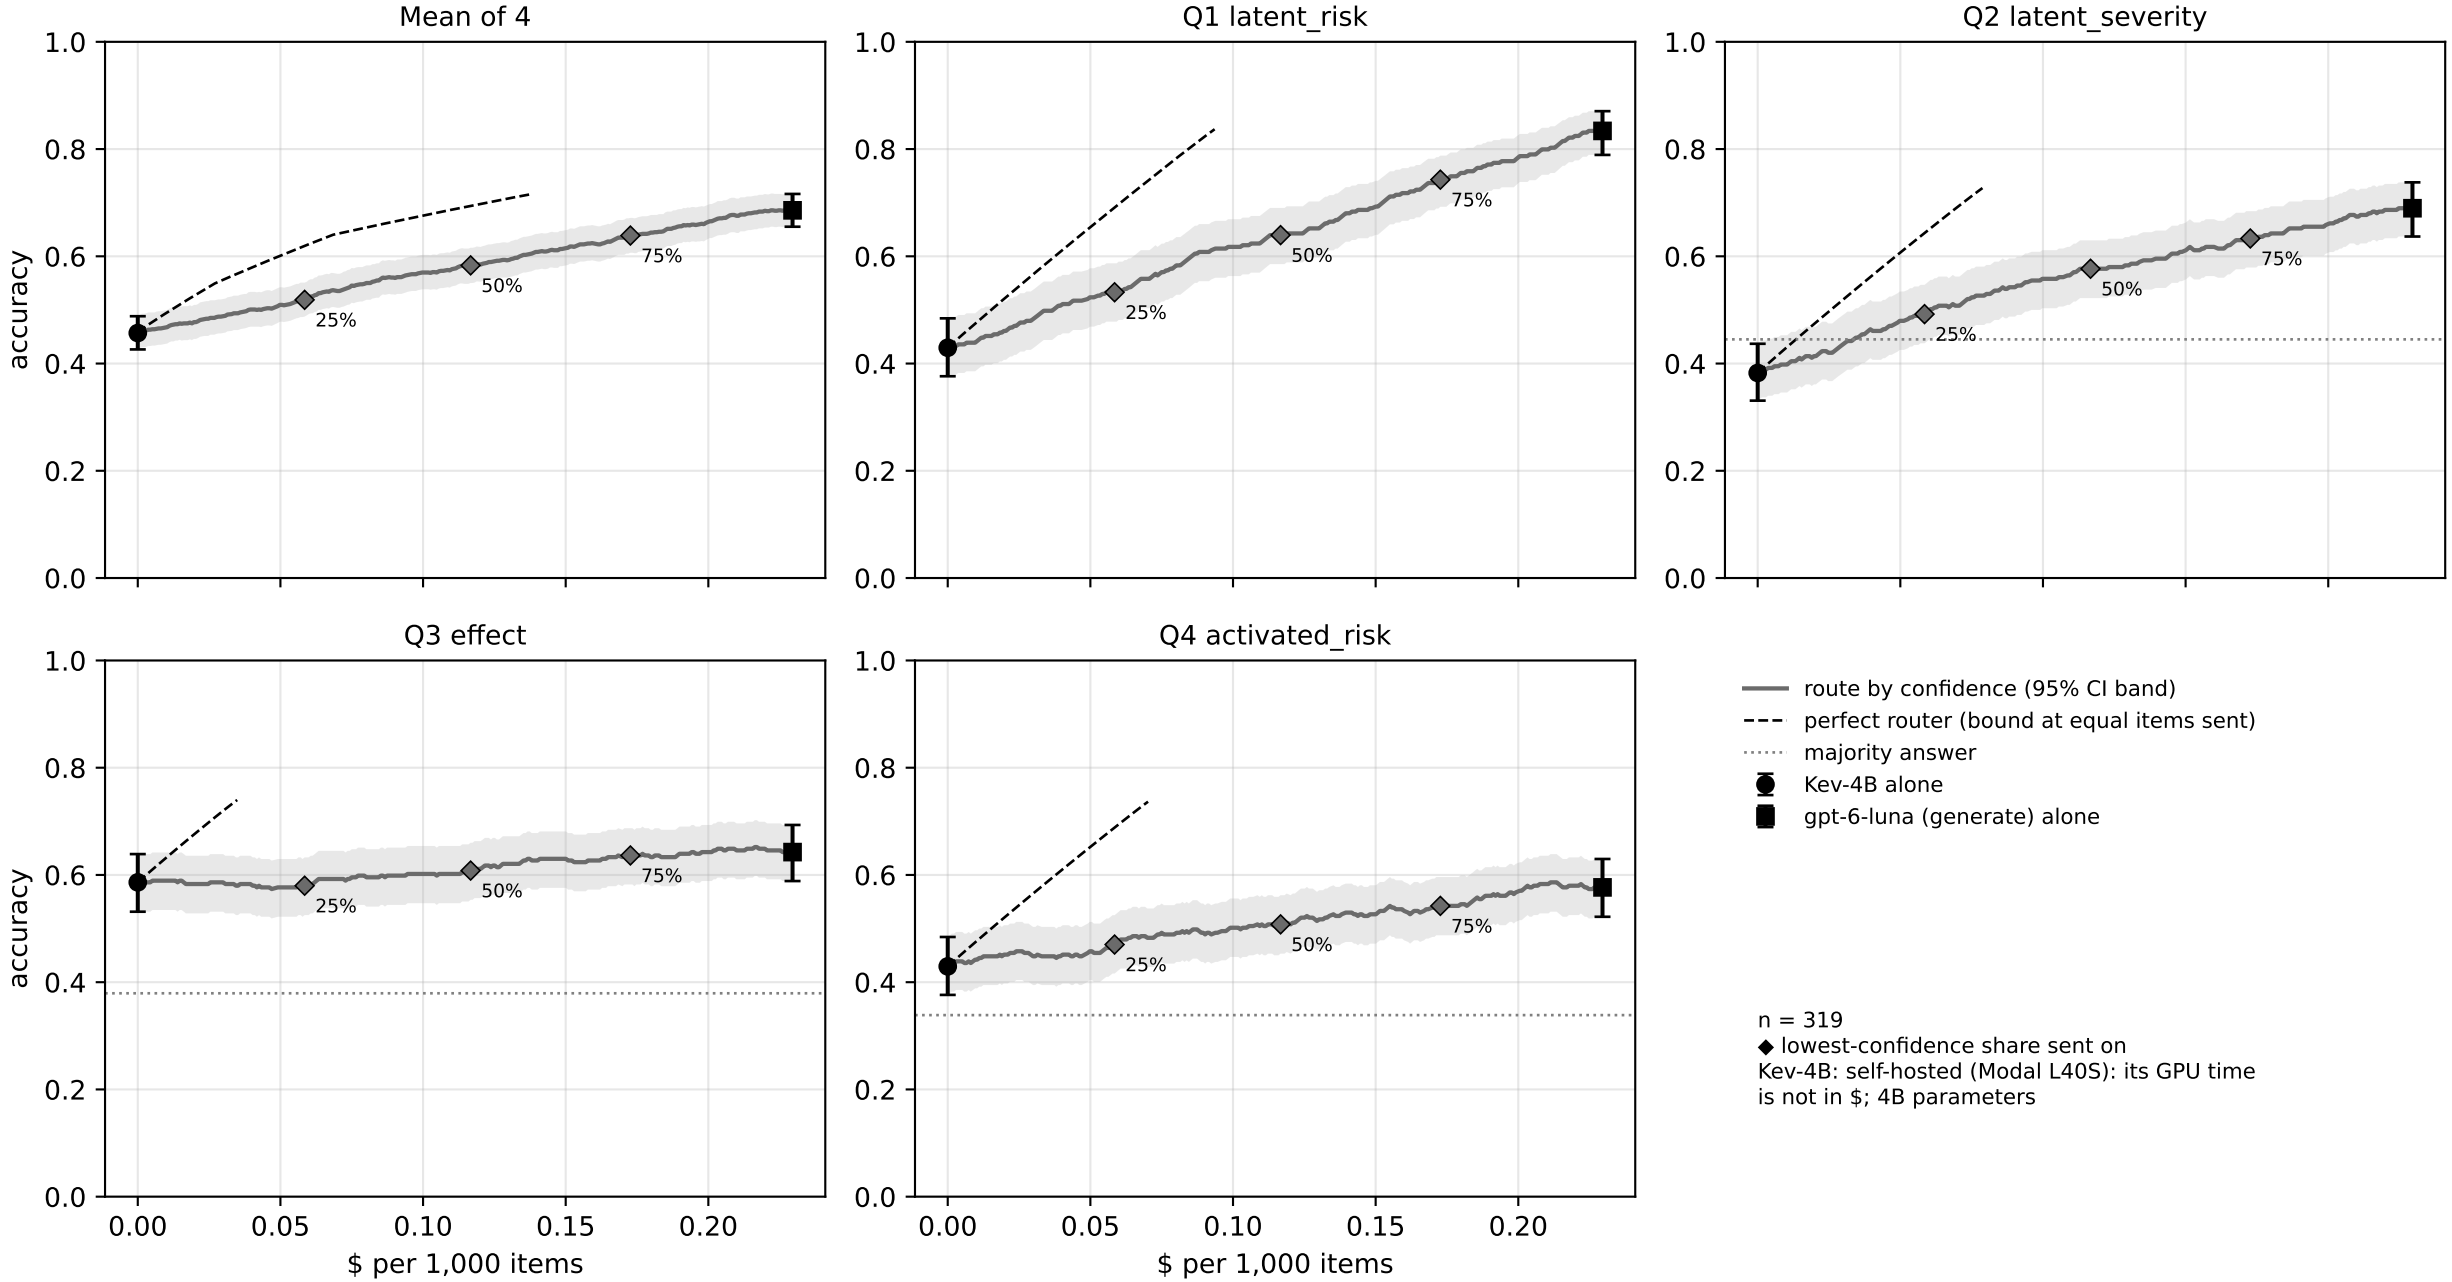
EXPLORATORY, post-hoc: escalation targets chosen after the Part 2 results were seen (not pre-registered)
Kev-4B first; low-confidence items to gpt-6-luna (generate)
Mean of 4
Q1 latent_risk
Q2 latent_severity
Q3 effect
Q4 activated_risk
accuracy
$ per 1,000 items
25%
50%
75%
route by confidence (95% CI band)
perfect router (bound at equal items sent)
majority answer
Kev-4B alone
gpt-6-luna (generate) alone
n = 319
◆ lowest-confidence share sent on
Kev-4B: self-hosted (Modal L40S): its GPU time
is not in $; 4B parameters