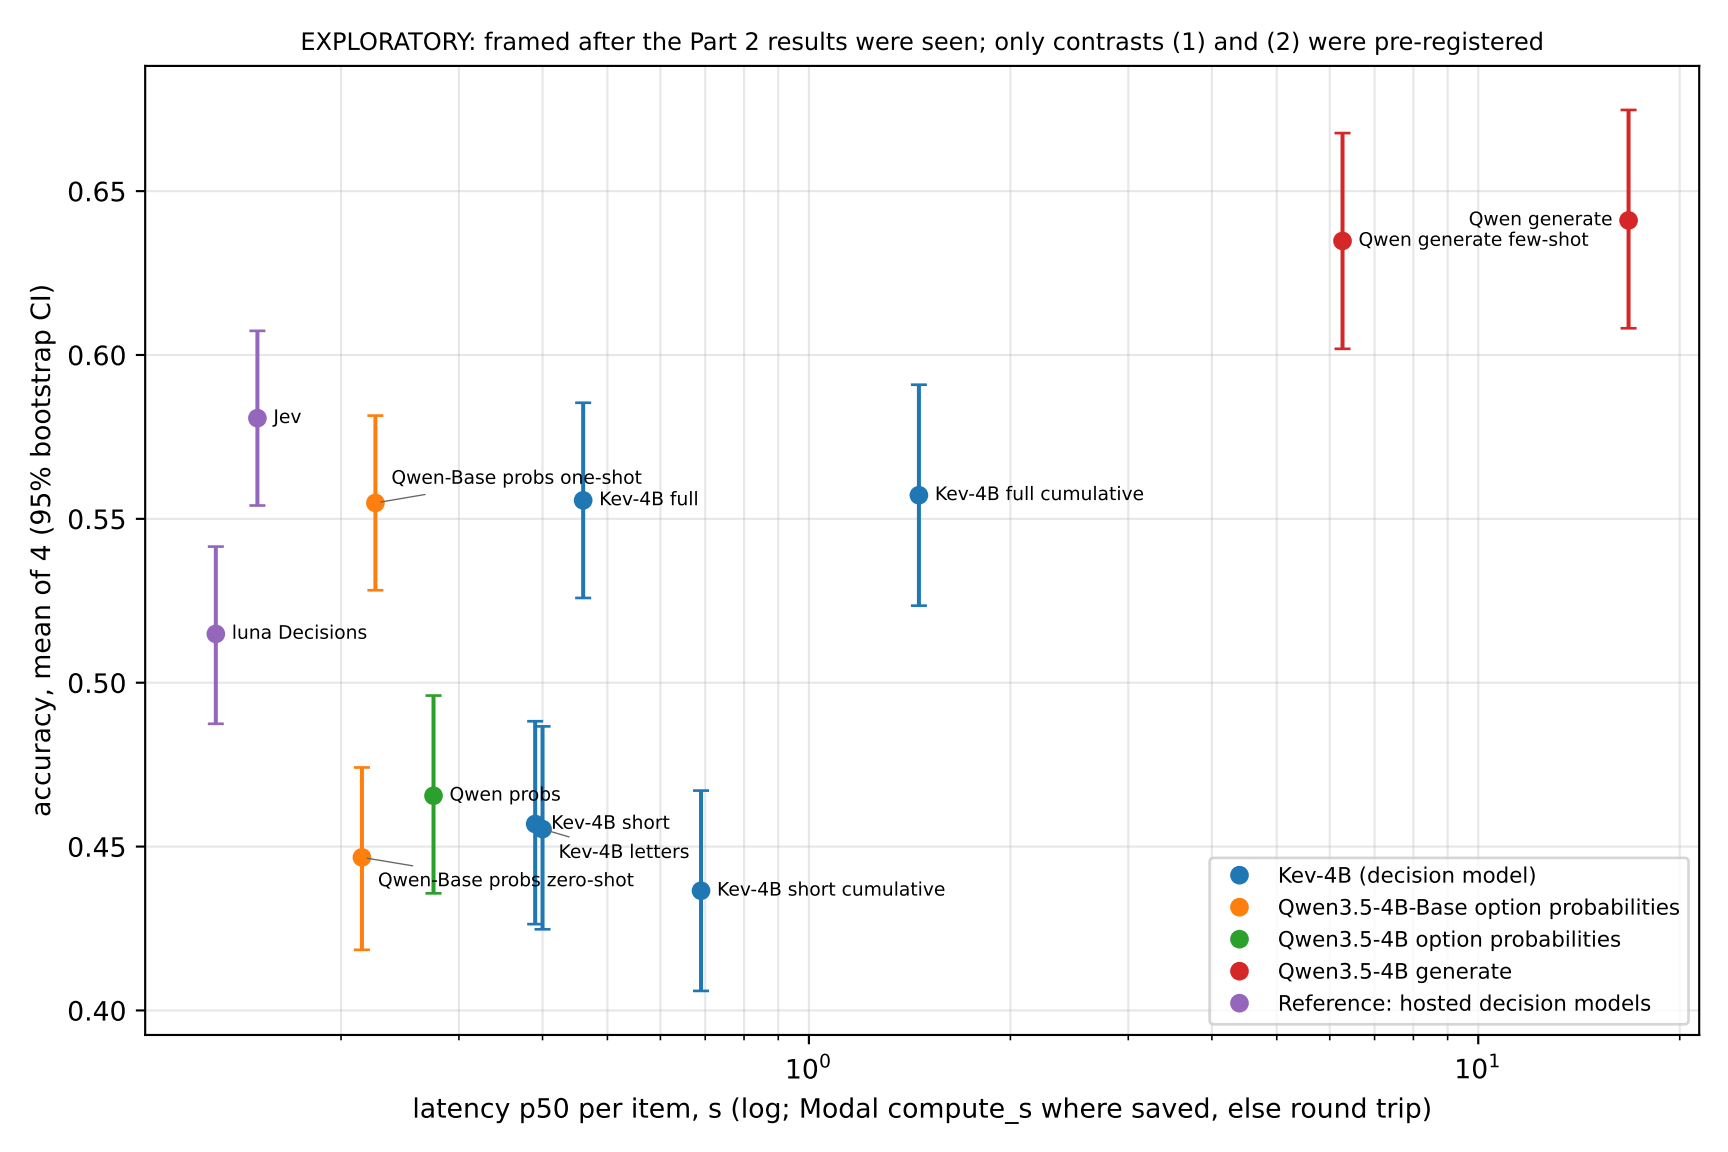
EXPLORATORY: framed after the Part 2 results were seen; only contrasts (1) and (2) were pre-registered
accuracy, mean of 4 (95% bootstrap CI)
0.65
0.60
0.55
0.50
0.45
0.40
10^0
10^1
latency p50 per item, s (log; Modal compute_s where saved, else round trip)
Qwen generate
Qwen generate few-shot
Jev
Qwen-Base probs one-shot
Kev-4B full
Kev-4B full cumulative
luna Decisions
Qwen probs
Kev-4B short
Kev-4B letters
Qwen-Base probs zero-shot
Kev-4B short cumulative
Kev-4B (decision model)
Qwen3.5-4B-Base option probabilities
Qwen3.5-4B option probabilities
Qwen3.5-4B generate
Reference: hosted decision models