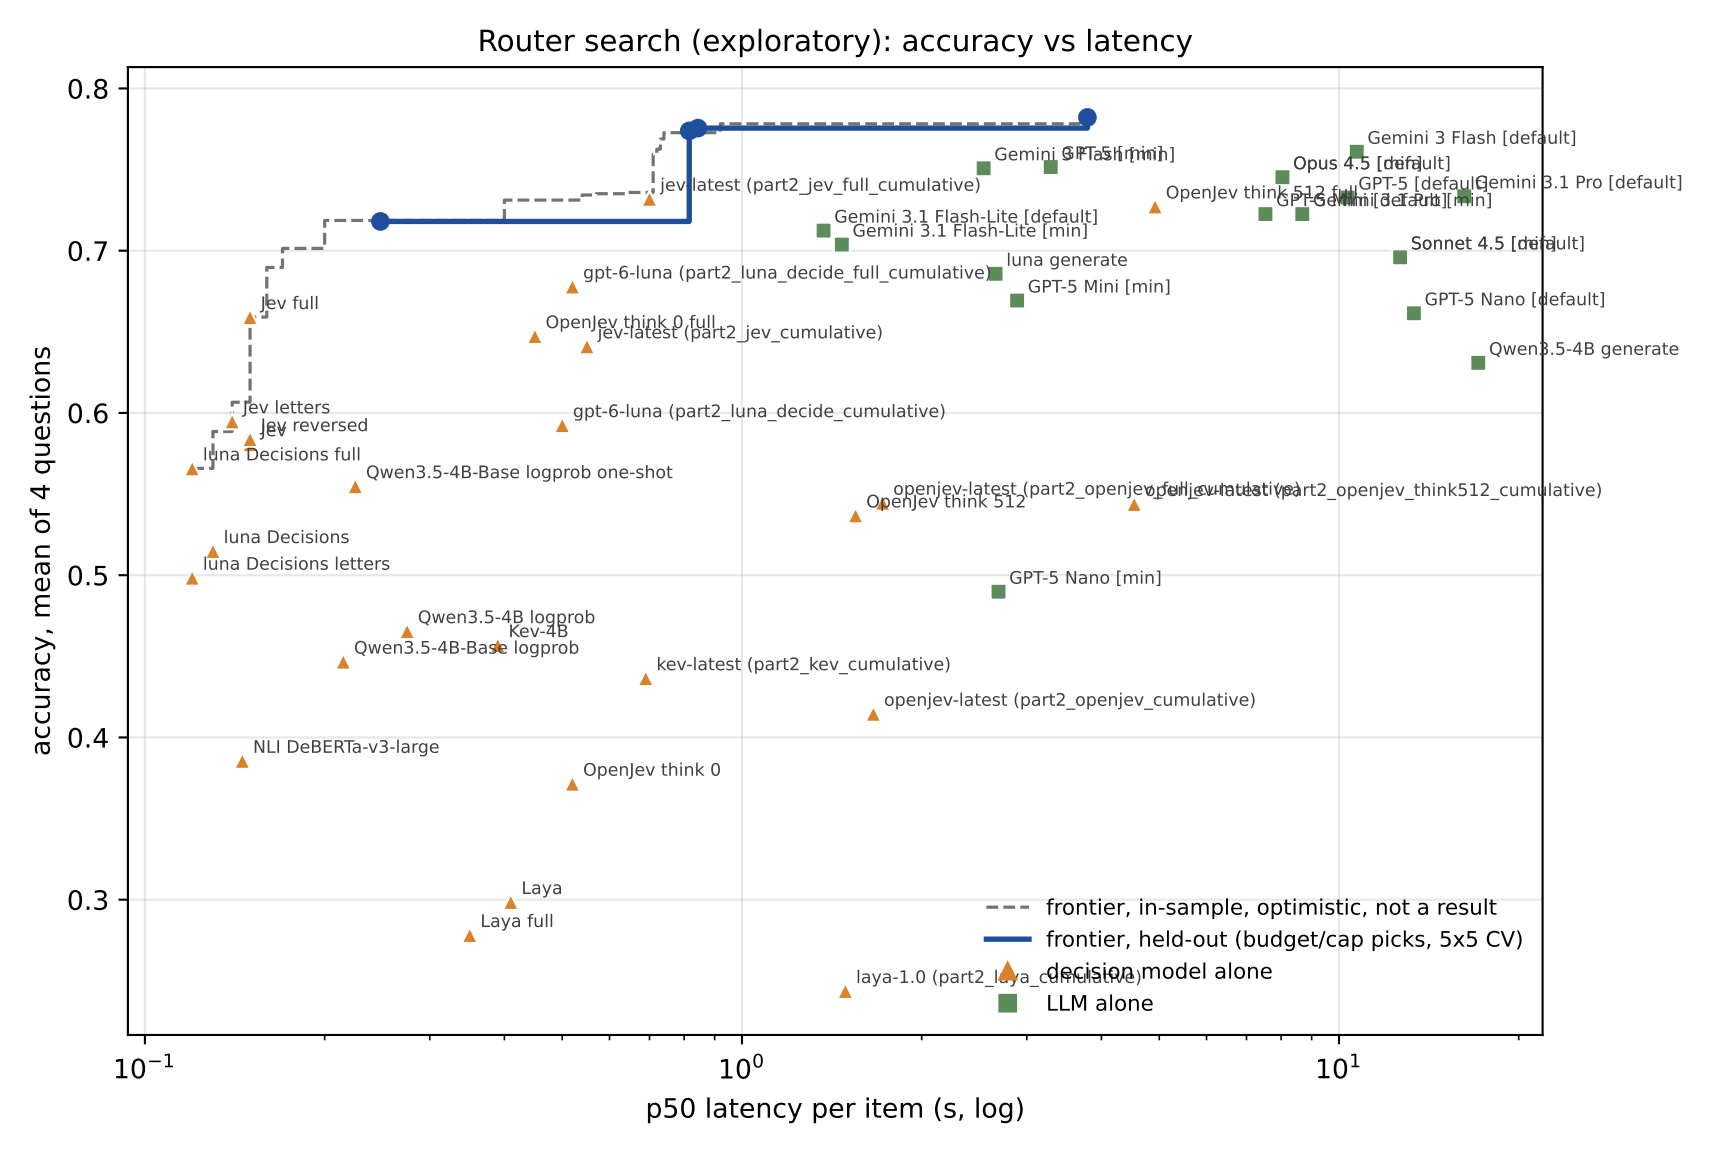
Router search (exploratory): accuracy vs latency
accuracy, mean of 4 questions
p50 latency per item (s, log)
0.8
0.7
0.6
0.5
0.4
0.3
10^{-1}
10^{0}
10^{1}
frontier, in-sample, optimistic, not a result
frontier, held-out (budget/cap picks, 5x5 CV)
decision model alone
LLM alone
jev-latest (part2_jev_full_cumulative)
Jev full
jev letters
Jev reversed
Jev
luna Decisions full
Qwen3.5-4B-Base logprob one-shot
luna Decisions
luna Decisions letters
Qwen3.5-4B logprob
Kev-4B
Qwen3.5-4B-Base logprob
NLI DeBERTa-v3-large
Laya
Laya full
gpt-6-luna (part2_luna_decide_full_cumulative)
OpenJev think 0 full
jev-latest (part2_jev_cumulative)
gpt-6-luna (part2_luna_decide_cumulative)
OpenJev think 0
kev-latest (part2_kev_cumulative)
openjev-latest (part2_openjev_cumulative)
laya-1.0 (part2_laya_cumulative)
OpenJev think 512
openjev-latest (part2_openjev_full_cumulative)
openjev-latest (part2_openjev_think512_cumulative)
OpenJev think 512 full
Gemini 3.1 Flash-Lite [default]
Gemini 3.1 Flash-Lite [min]
luna generate
GPT-5 Mini [min]
GPT-5 Nano [min]
Gemini 3 Flash [default]
Gemini 3 Flash [min]
Opus 4.5 [default]
GPT-5 [default]
Gemini 3.1 Pro [default]
Sonnet 4.5 [default]
GPT-5 Nano [default]
Qwen3.5-4B generate