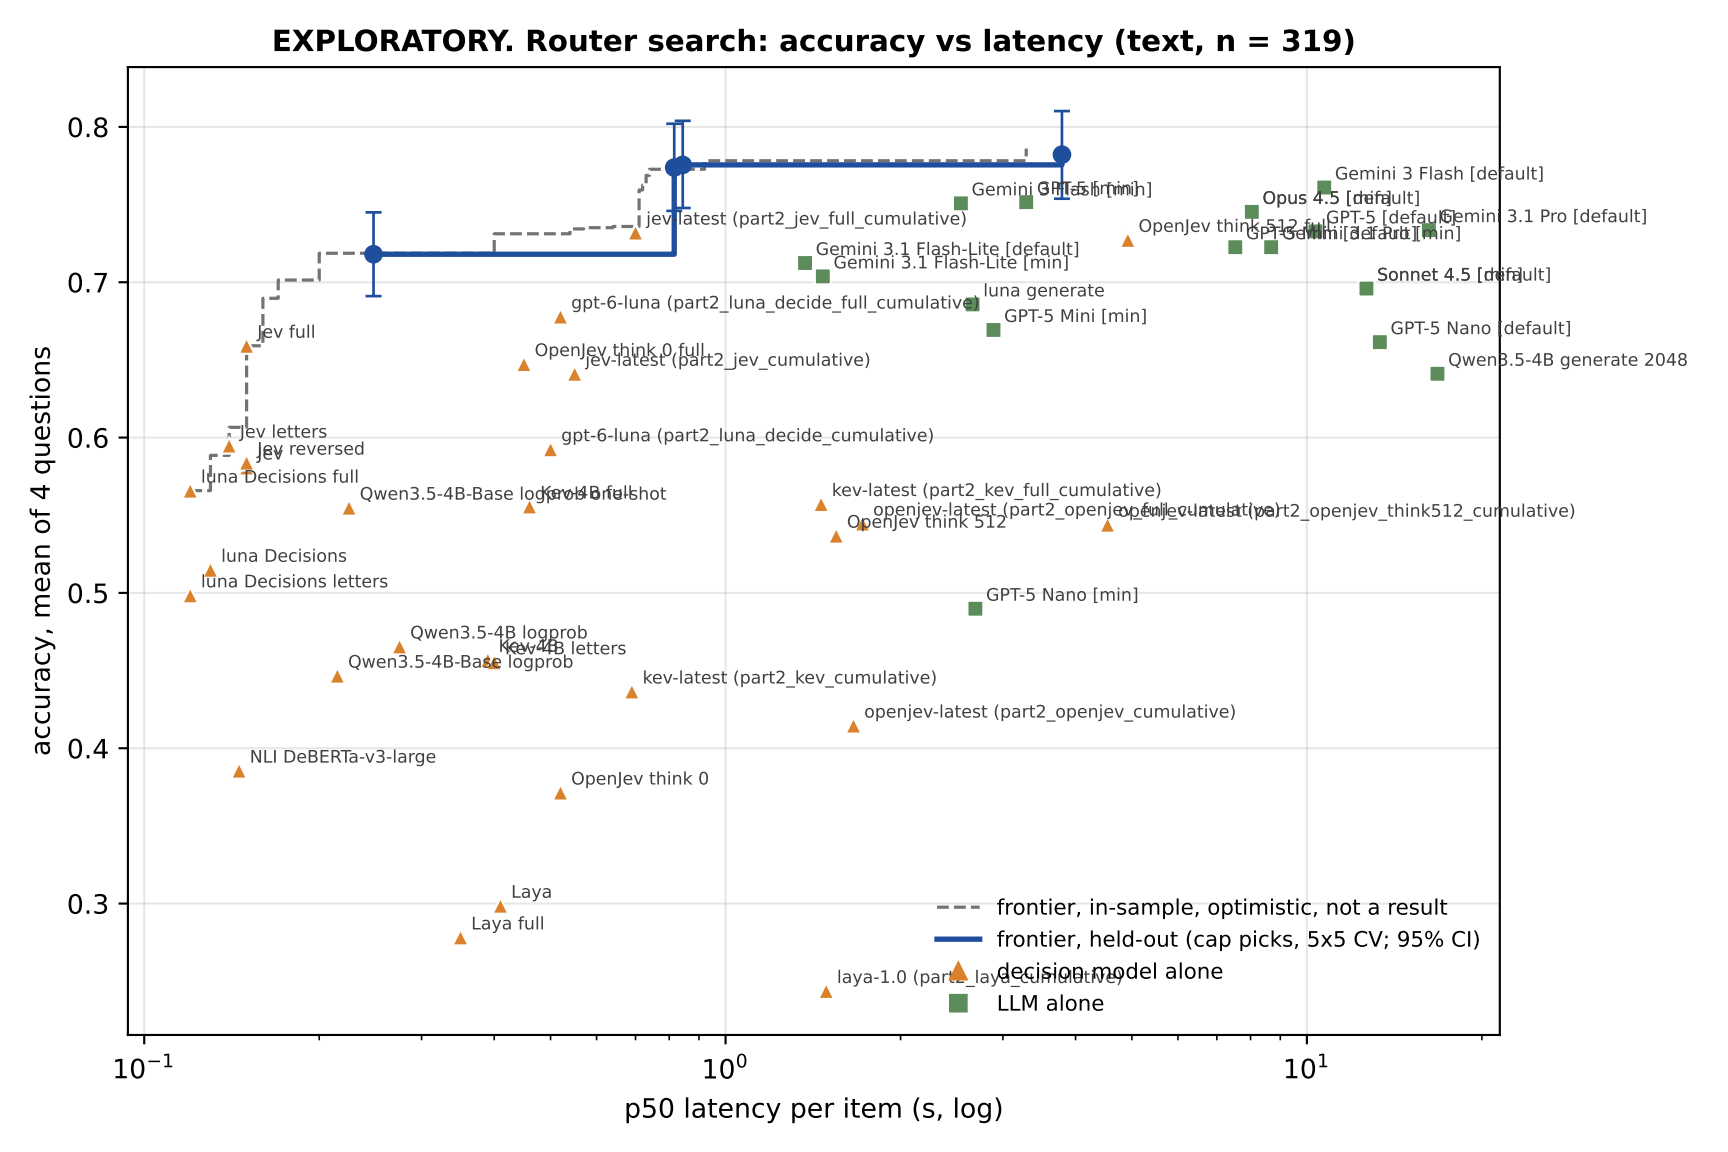
EXPLORATORY. Router search: accuracy vs latency (text, n = 319)
accuracy, mean of 4 questions
p50 latency per item (s, log)
0.8
0.7
0.6
0.5
0.4
0.3
10^{-1}
10^{0}
10^{1}
Gemini 3 Flash [default]
Gemini 3.1 Pro [default]
Sonnet 4.5 [default]
GPT-5 Nano [default]
Qwen3.5-4B generate 2048
GPT-5 Mini [min]
luna generate
Gemini 3.1 Flash-Lite [default]
Gemini 3.1 Flash-Lite [min]
jev latest (part2_jev_full_cumulative)
gpt-6-luna (part2_luna_decide_full_cumulative)
jev full
OpenJev think 0 full
jev-latest (part2_jev_cumulative)
jev letters
jev reversed
jev
luna Decisions full
gpt-6-luna (part2_luna_decide_cumulative)
Qwen3.5-4B-Base logprob 0-shot
kev-latest (part2_kev_full_cumulative)
OpenJev think 512
GPT-5 Nano [min]
luna Decisions
luna Decisions letters
Qwen3.5-4B logprob
Qwen3.5-4B-Base logprob
kev-latest (part2_kev_cumulative)
openjev-latest (part2_openjev_cumulative)
NLI DeBERTa-v3-large
OpenJev think 0
Laya
Laya full
frontier, in-sample, optimistic, not a result
frontier, held-out (cap picks, 5x5 CV; 95% CI)
decision model alone
LLM alone