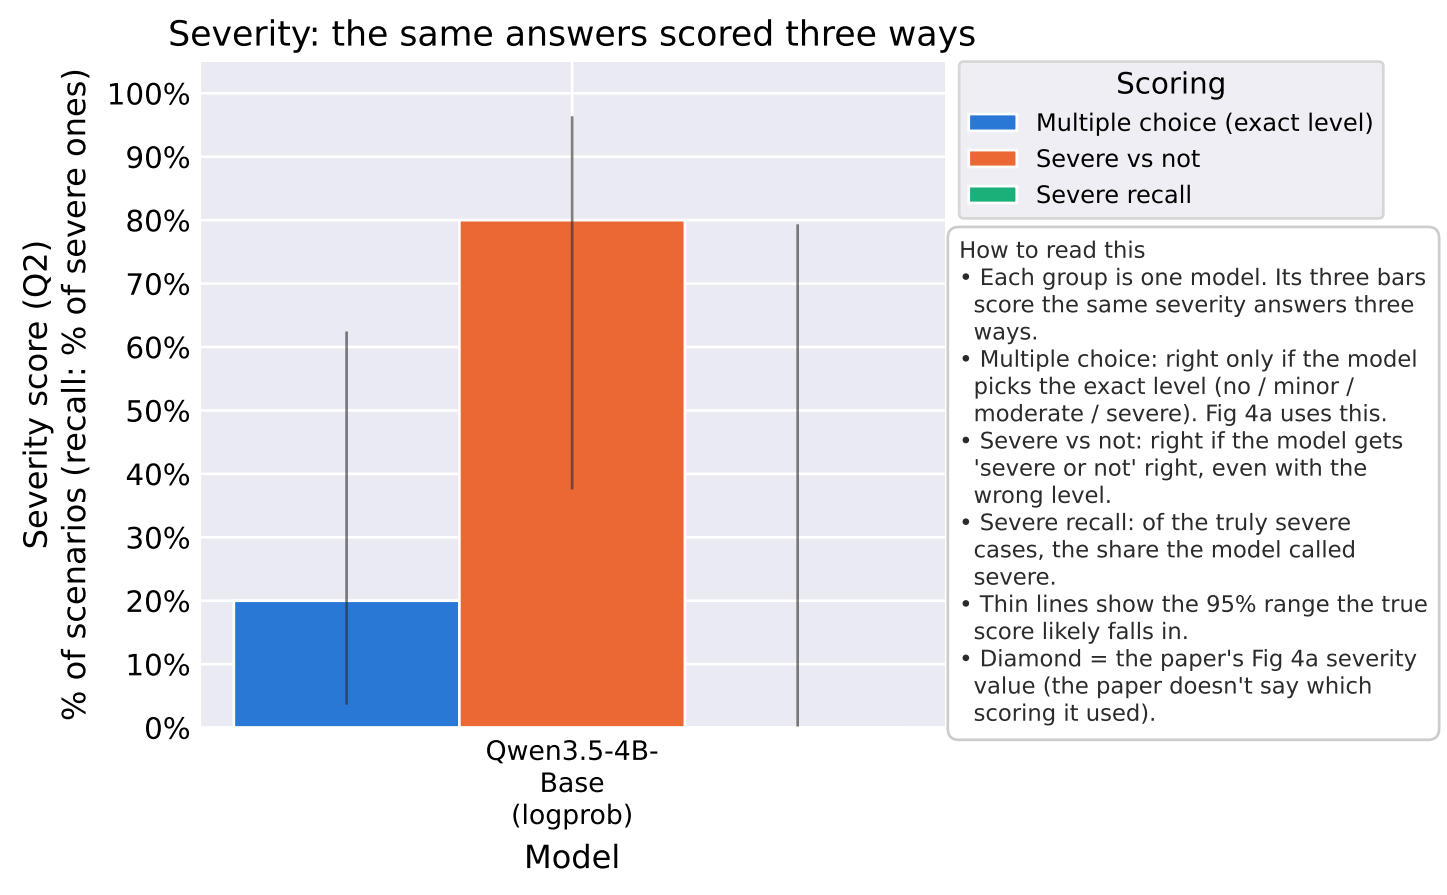
Severity: the same answers scored three ways
Severity score (Q2)
% of scenarios (recall: % of severe ones)
100%
90%
80%
70%
60%
50%
40%
30%
20%
10%
0%
Qwen3.5-4B-Base (logprob)
Model
Scoring
Multiple choice (exact level)
Severe vs not
Severe recall
How to read this
• Each group is one model. Its three bars score the same severity answers three ways.
• Multiple choice: right only if the model picks the exact level (no / minor / moderate / severe). Fig 4a uses this.
• Severe vs not: right if the model gets 'severe or not' right, even with the wrong level.
• Severe recall: of the truly severe cases, the share the model called severe.
• Thin lines show the 95% range the true score likely falls in.
• Diamond = the paper's Fig 4a severity value (the paper doesn't say which scoring it used).