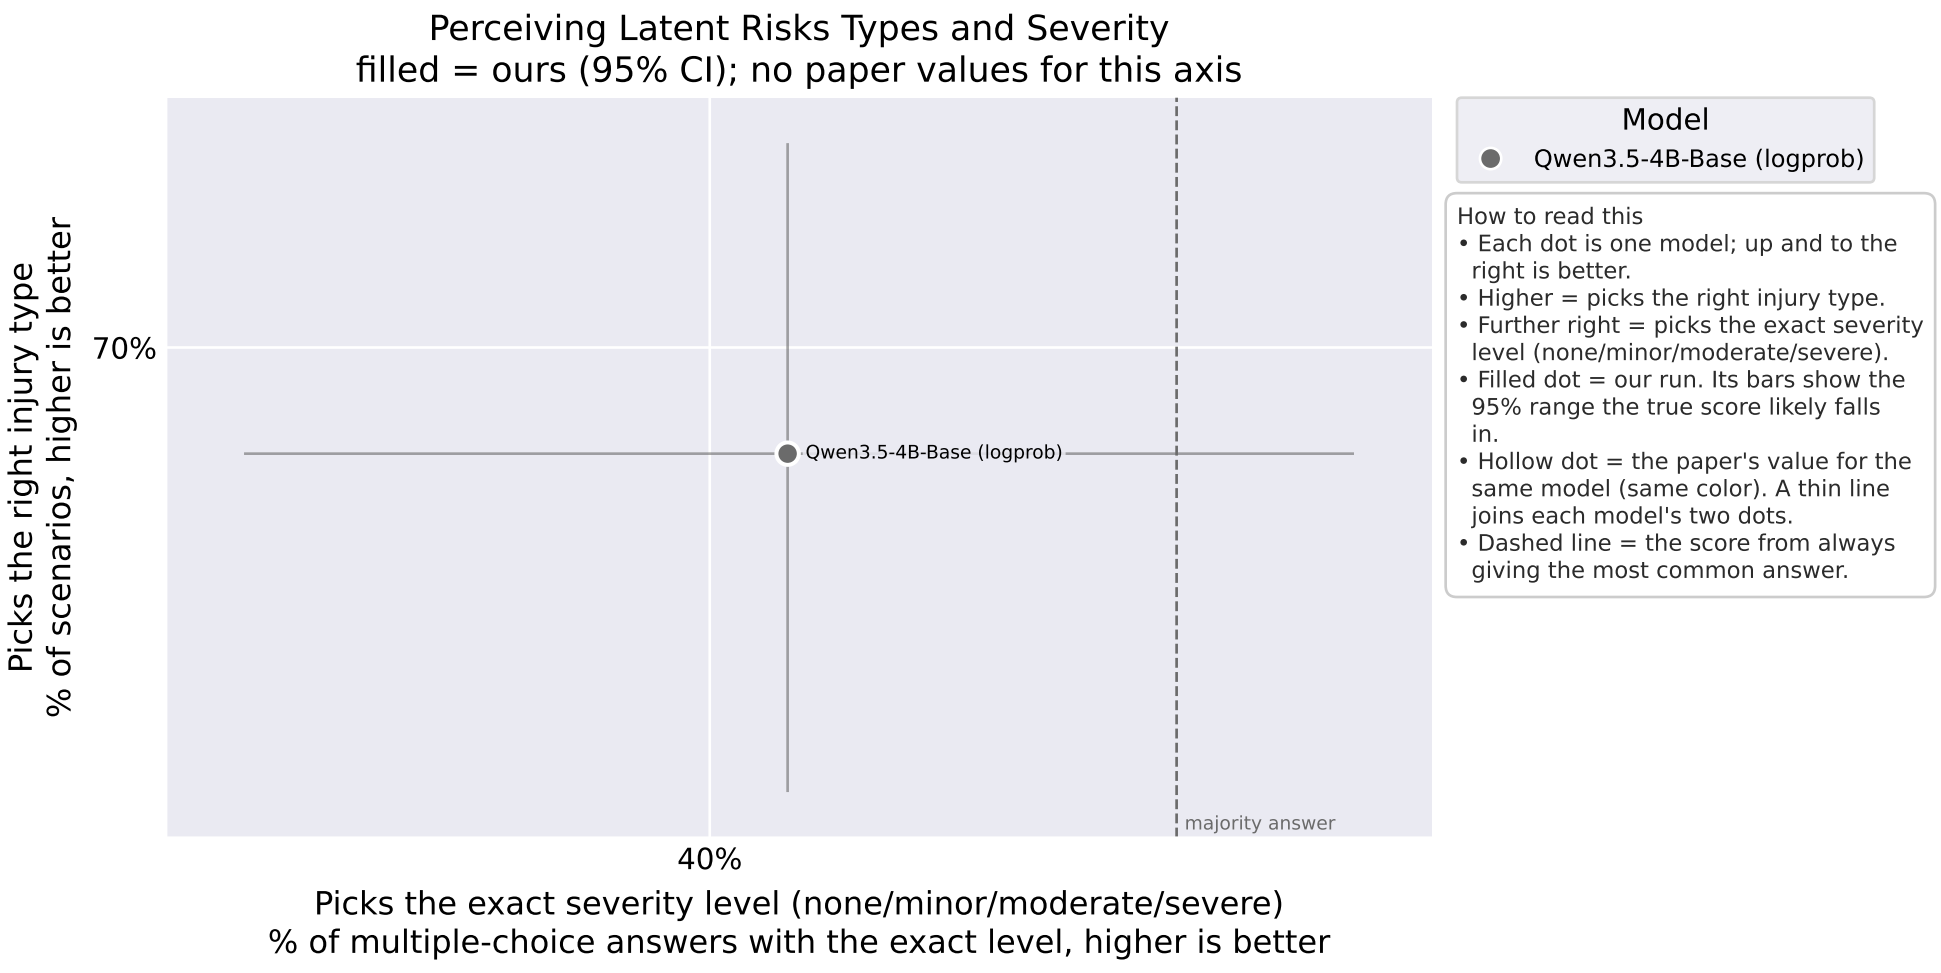
Perceiving Latent Risks Types and Severity
filled = ours (95% CI); no paper values for this axis
Picks the right injury type
% of scenarios, higher is better
70%
40%
Qwen3.5-4B-Base (logprob)
majority answer
Picks the exact severity level (none/minor/moderate/severe)
% of multiple-choice answers with the exact level, higher is better
Model
Qwen3.5-4B-Base (logprob)
How to read this
• Each dot is one model; up and to the right is better.
• Higher = picks the right injury type.
• Further right = picks the exact severity level (none/minor/moderate/severe).
• Filled dot = our run. Its bars show the 95% range the true score likely falls in.
• Hollow dot = the paper's value for the same model (same color). A thin line joins each model's two dots.
• Dashed line = the score from always giving the most common answer.