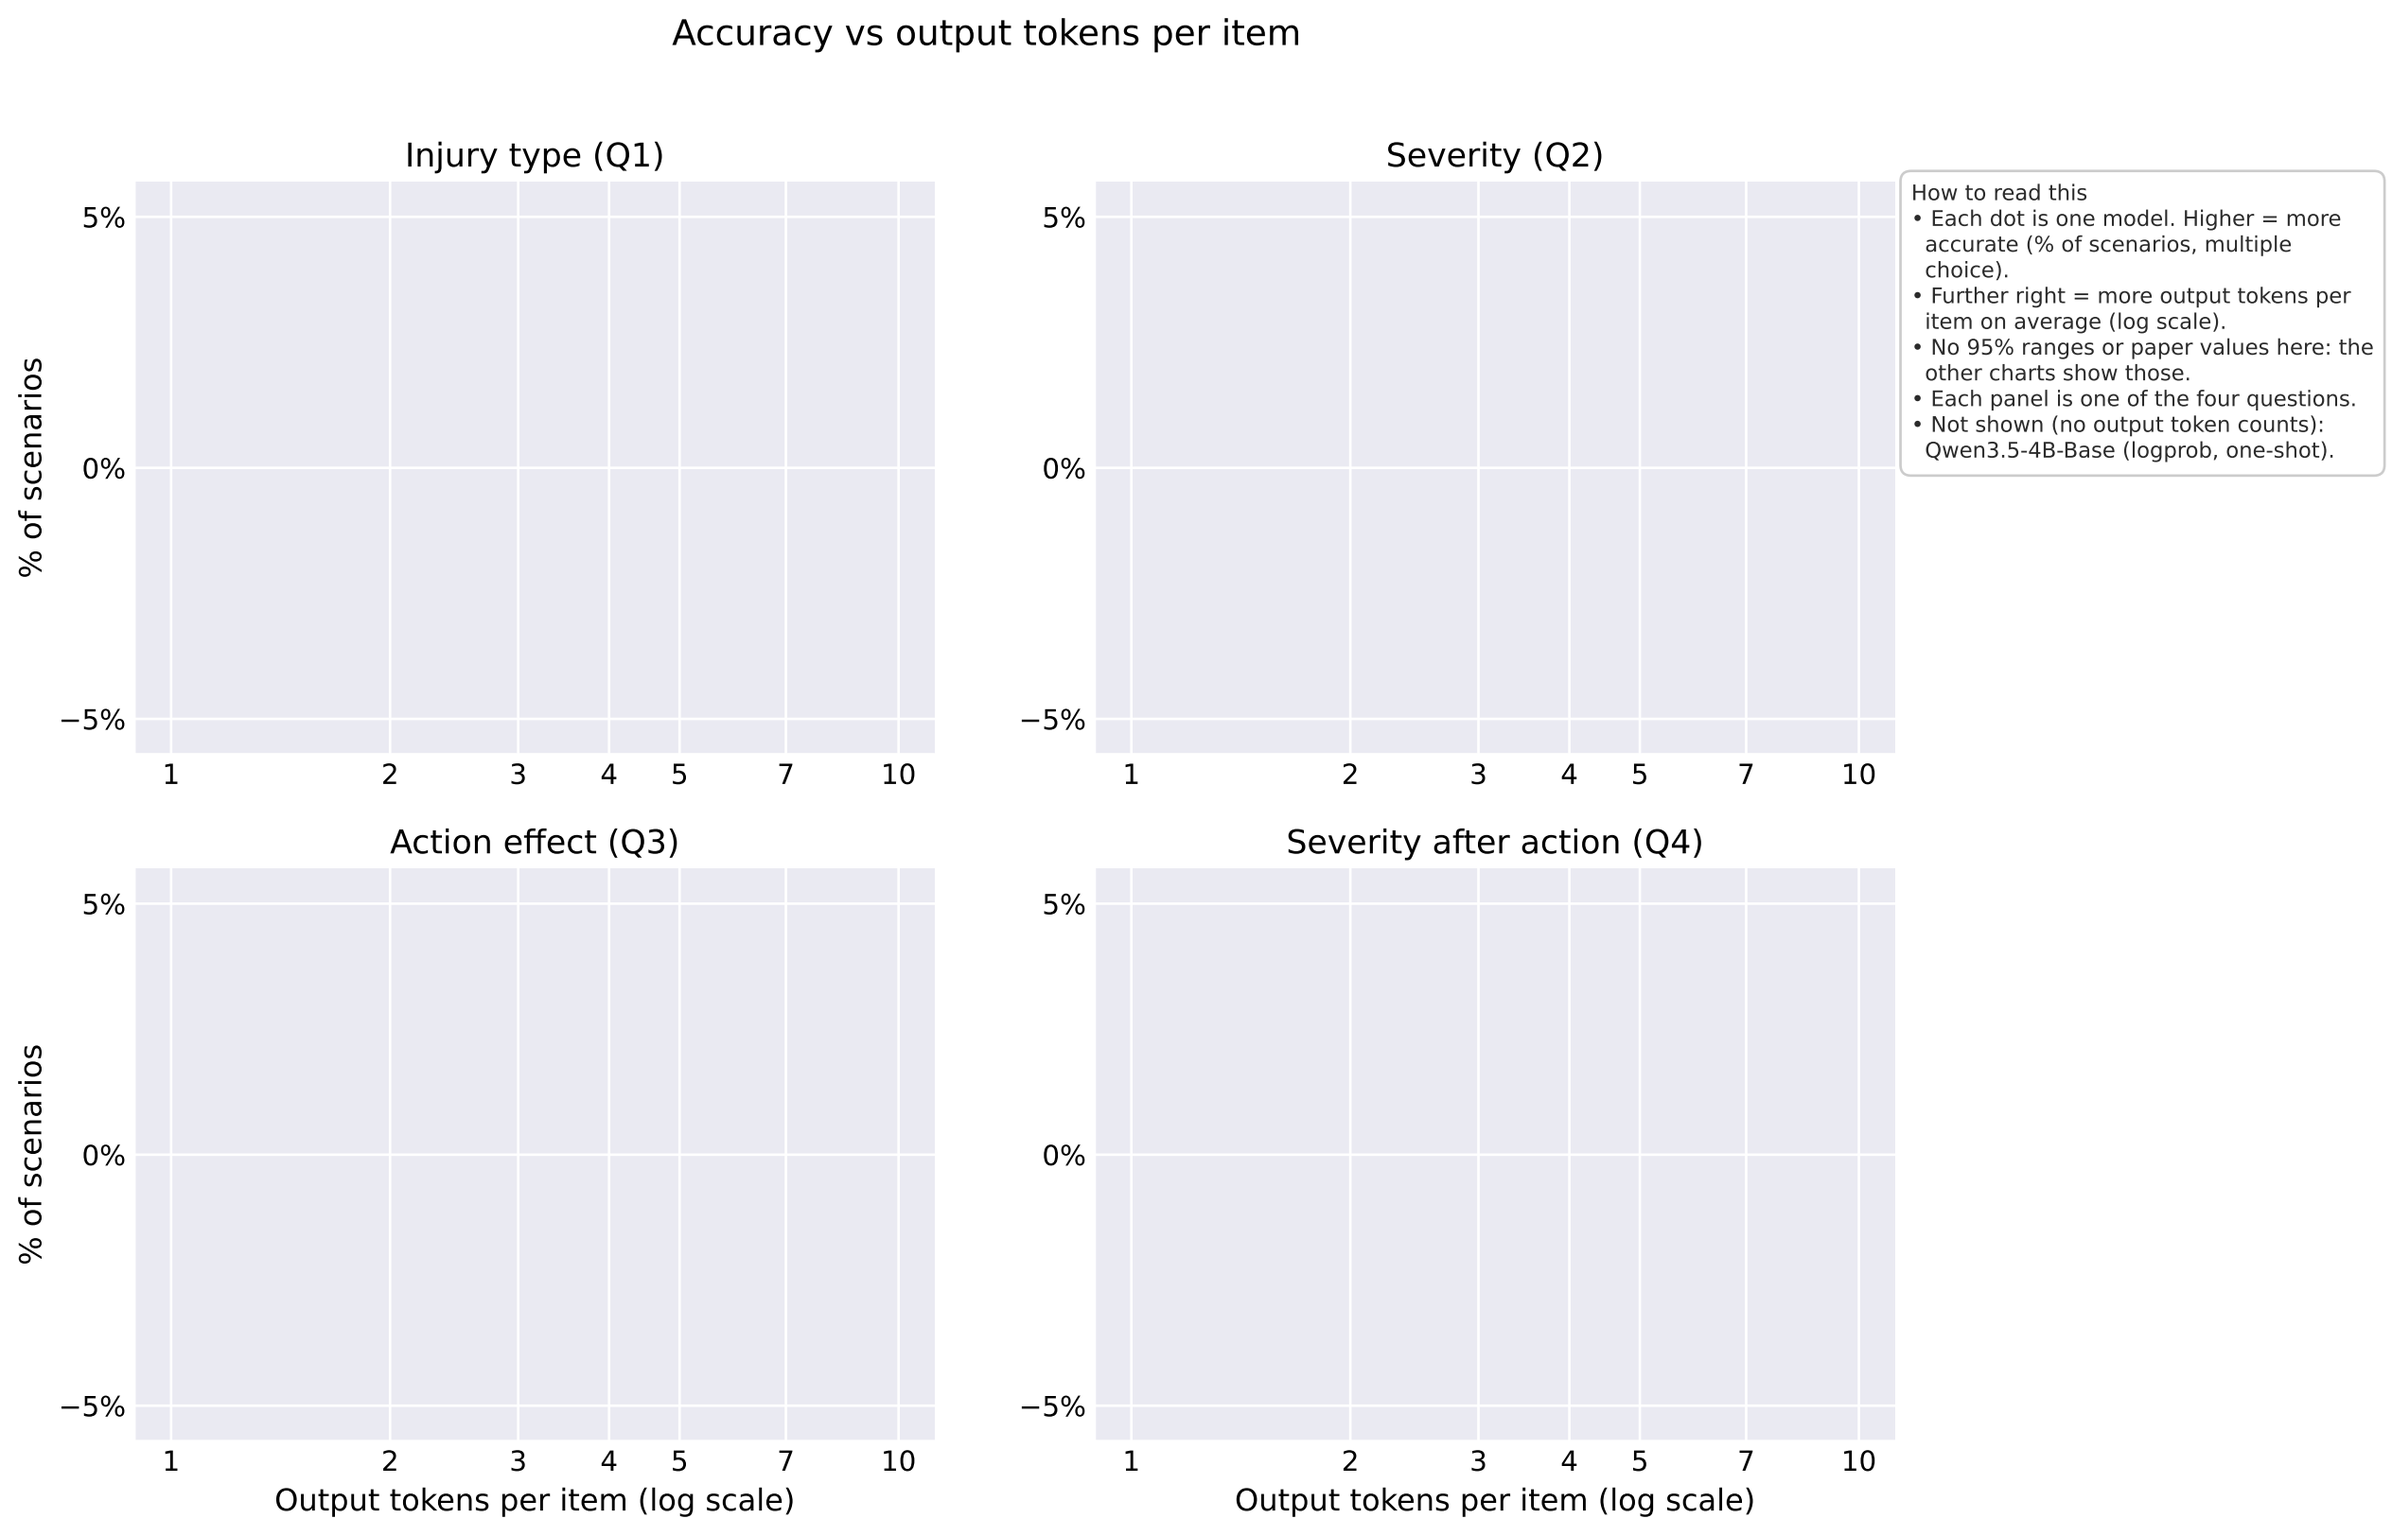
Accuracy vs output tokens per item
Injury type (Q1)
Severity (Q2)
Action effect (Q3)
Severity after action (Q4)
% of scenarios
5%
0%
−5%
1
2
3
4
5
7
10
Output tokens per item (log scale)
How to read this
• Each dot is one model. Higher = more accurate (% of scenarios, multiple choice).
• Further right = more output tokens per item on average (log scale).
• No 95% ranges or paper values here: the other charts show those.
• Each panel is one of the four questions.
• Not shown (no output token counts): Qwen3.5-4B-Base (logprob, one-shot).