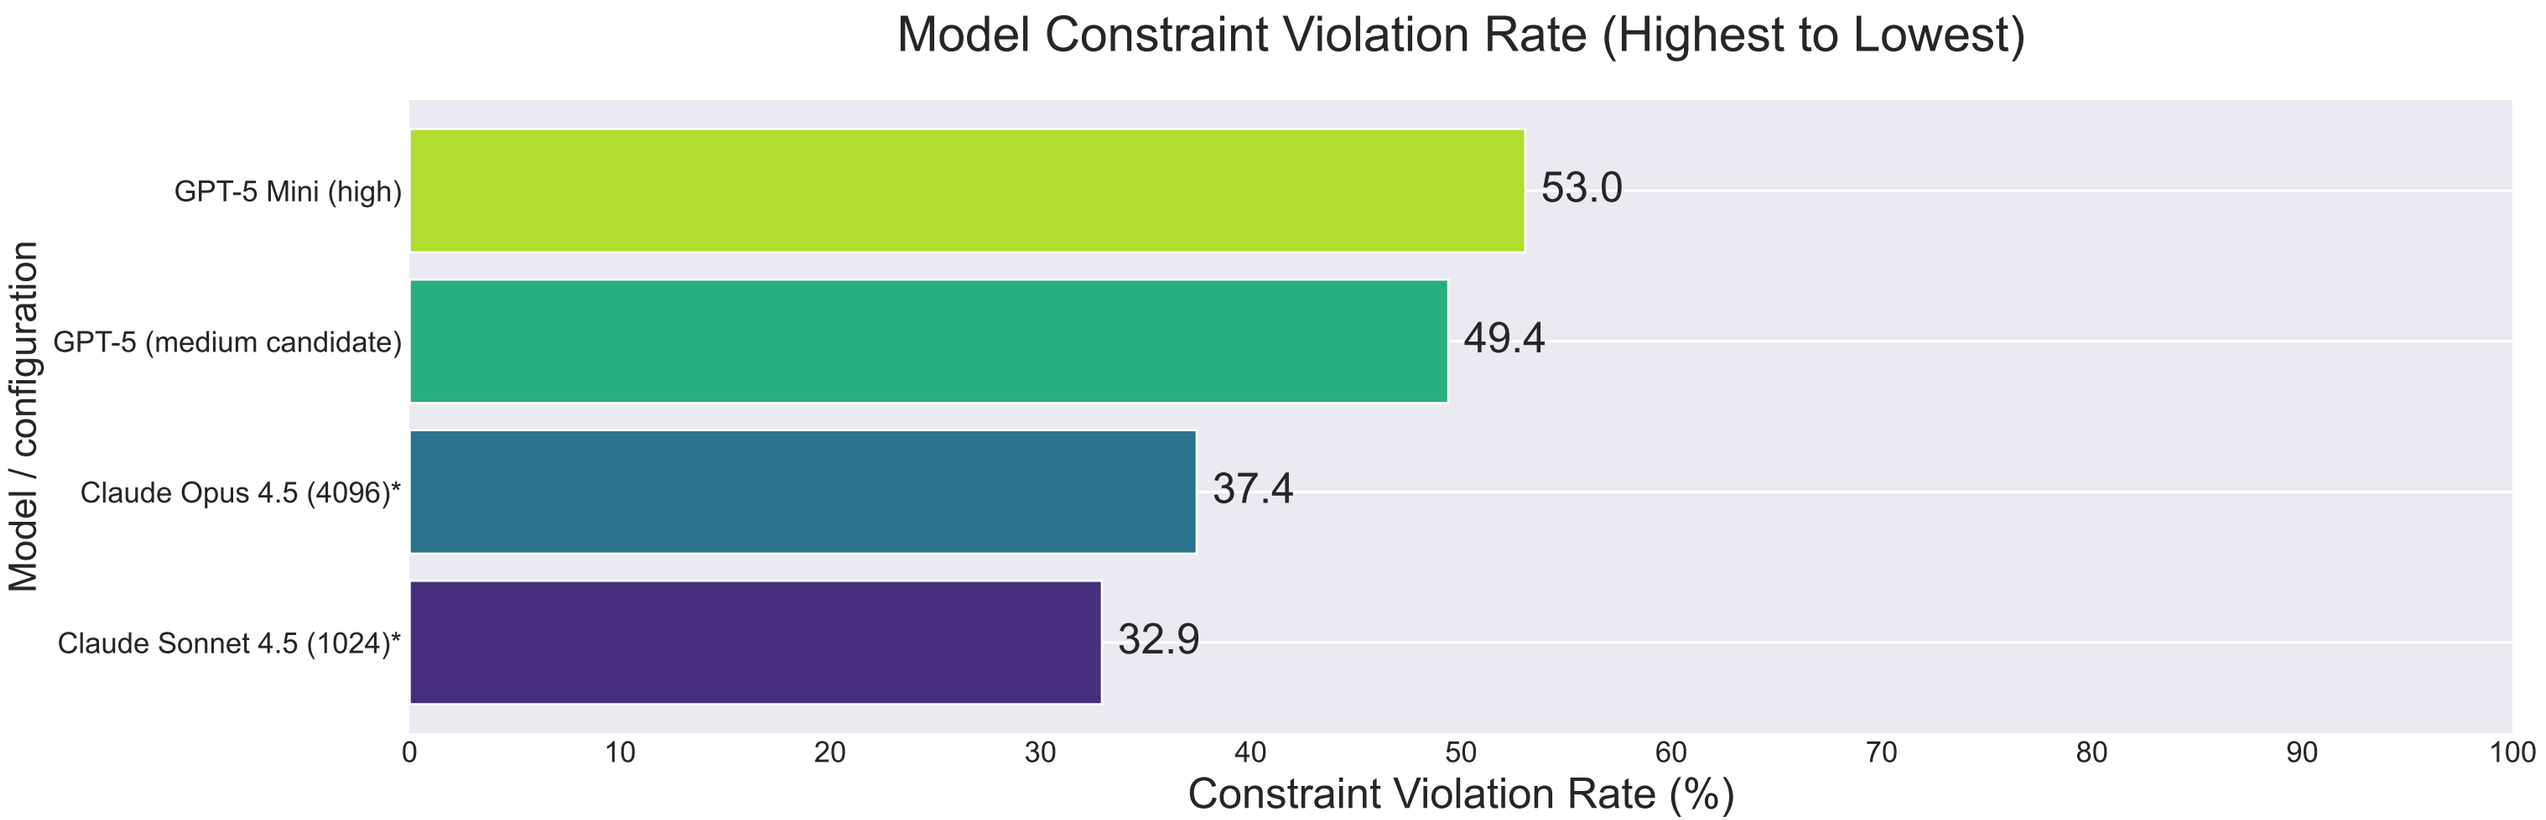

164 examples per full run; rates exclude malformed responses. Lower is better. Settings and counts are in the CSV.

* Claude 4.5 substitutes: exact Sonnet 4 / Opus 4.1 API IDs returned 404.

Settings fixed before the experiment; same literal scorer. Original paper inference/aggregation settings remain unverified.