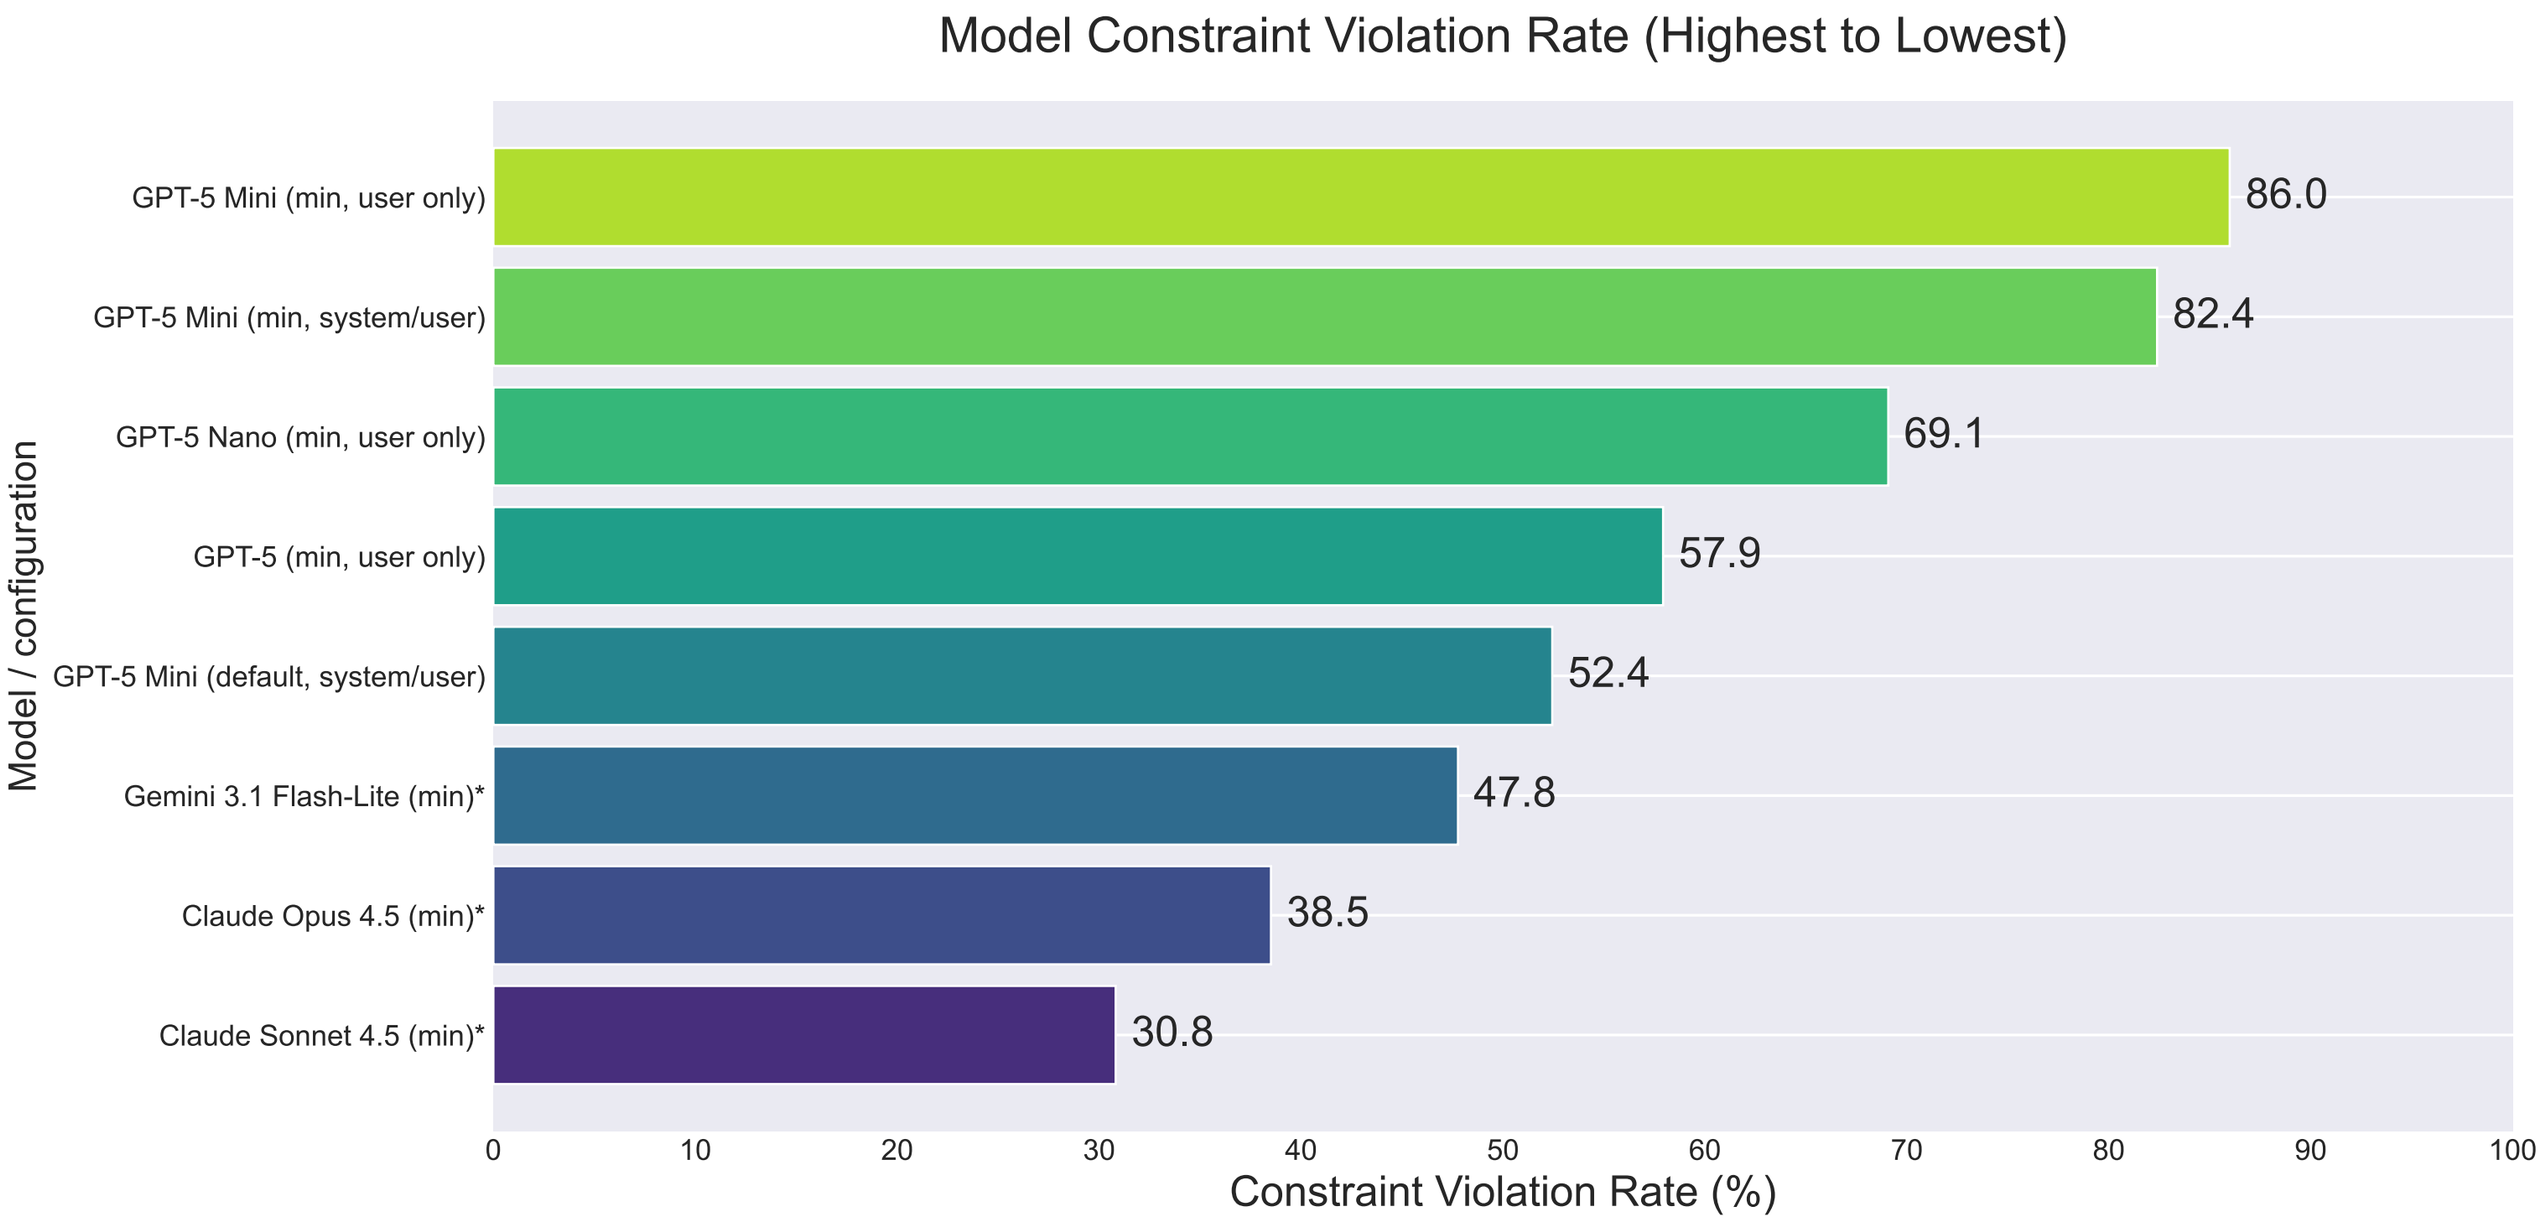

164 examples per full run; rates exclude malformed responses. Lower is better. Settings and counts are in the CSV.

* Newer substitutes for the paper's Gemini 2.5 / Claude 4 models. Mixed thinking/role settings are labeled.

Gemini Flash Preview (22/164 responses) and Pro (0/164) omitted. These are reproduction attempts, not exact paper results.